Section A: Personal Workflow Audit

#	Recurring Weekly Task	Classification	One-Line Rationale
1	**Handcrafting organic shampoo batches**	**Just Me**	Physical, artisan-level compounding done in small batches by my mother that requires tactile quality control.
2	**Negotiating raw ingredient prices with local vendors**	**Just Me**	Requires building personal trust, navigating local wholesale markets, and real-time human negotiation.
3	**Testing and debugging physical hardware circuits**	**Just Me**	Requires physical laboratory instrumentation, multimeters, and manual wiring that AI cannot physically interact with.
4	**Writing complex SQL & DuckDB queries**	**Collaborate with AI**	I provide the database schema and query intent; the AI rapidly drafts the initial syntax, which I then optimize for performance.
5	**Drafting social copy & catalog descriptions for @Khaam.essentials**	**Collaborate with AI**	I provide the product benefits; the AI refines the tone and structures the copy, which I then localize for our target audience.
6	**Translating numerical analysis algorithms into Python scripts**	**Collaborate with AI**	Great for translating step-by-step mathematical logic (like Jacobi or Gauss-Seidel) into structured, clean code.

7	**Synthesizing engineering handbook chapters into review sheets**	**Delegate to AI with Review**	I upload the raw PDF chapter; the AI extracts key formulas and definitions, which I audit against my lecture notes.
8	**Reviewing documentation on rule-based recommendation scoring**	**Delegate to AI with Review**	The AI summarizes dense API/framework docs so I can quickly extract implementation patterns for my internship tasks.
9	**Triaging basic customer FAQs on Instagram & WhatsApp**	**Delegate to AI with Review**	Standard questions about shipping, product availability, or price are drafted by AI templates, requiring only my manual click to send.
10	**Formatting, linting, and cleaning local Python code bases**	**Fully Automate**	Easily managed using automated Git hooks, Black, and Flake8 formatters running locally on file save.
11	**Syncing raw inventory lists with spreadsheet stock counts**	**Fully Automate**	Handled by simple Google Sheets script triggers that auto-update stock levels whenever a WhatsApp order form is logged.
12	**Generating mock technical interview questions**	**Just Me**	Helps me practice core engineering and discrete modeling concepts through interactive chat sessions.

Section B: Setup Evidence & Custom Instructions

Paste this description into your document, and then **insert your screenshot** directly below it:

- **Claude.ai / ChatGPT Free Accounts:** Confirmed active.
- **Anthropic Academy:** Enrolled in *Claude 101* (Module 1 completed).

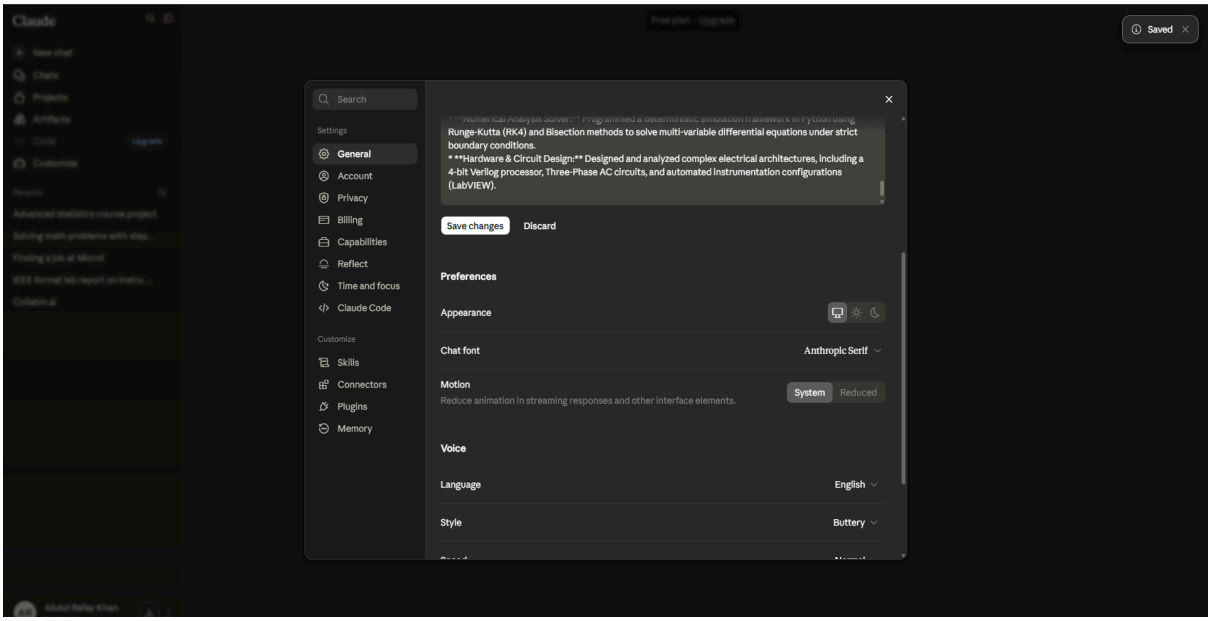

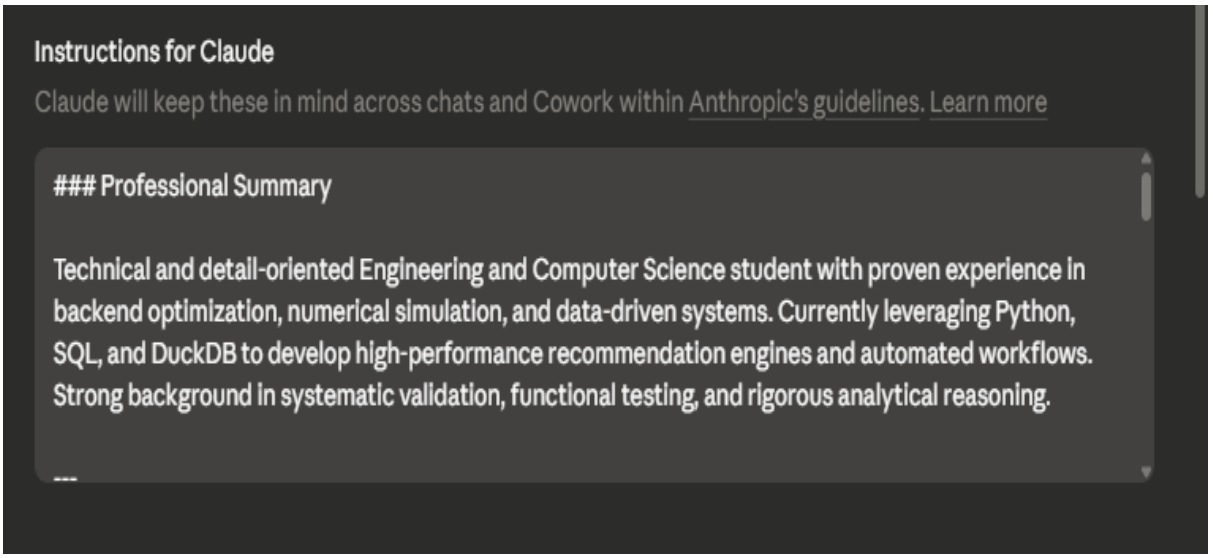

Section C: Selected Tasks & Success Metrics

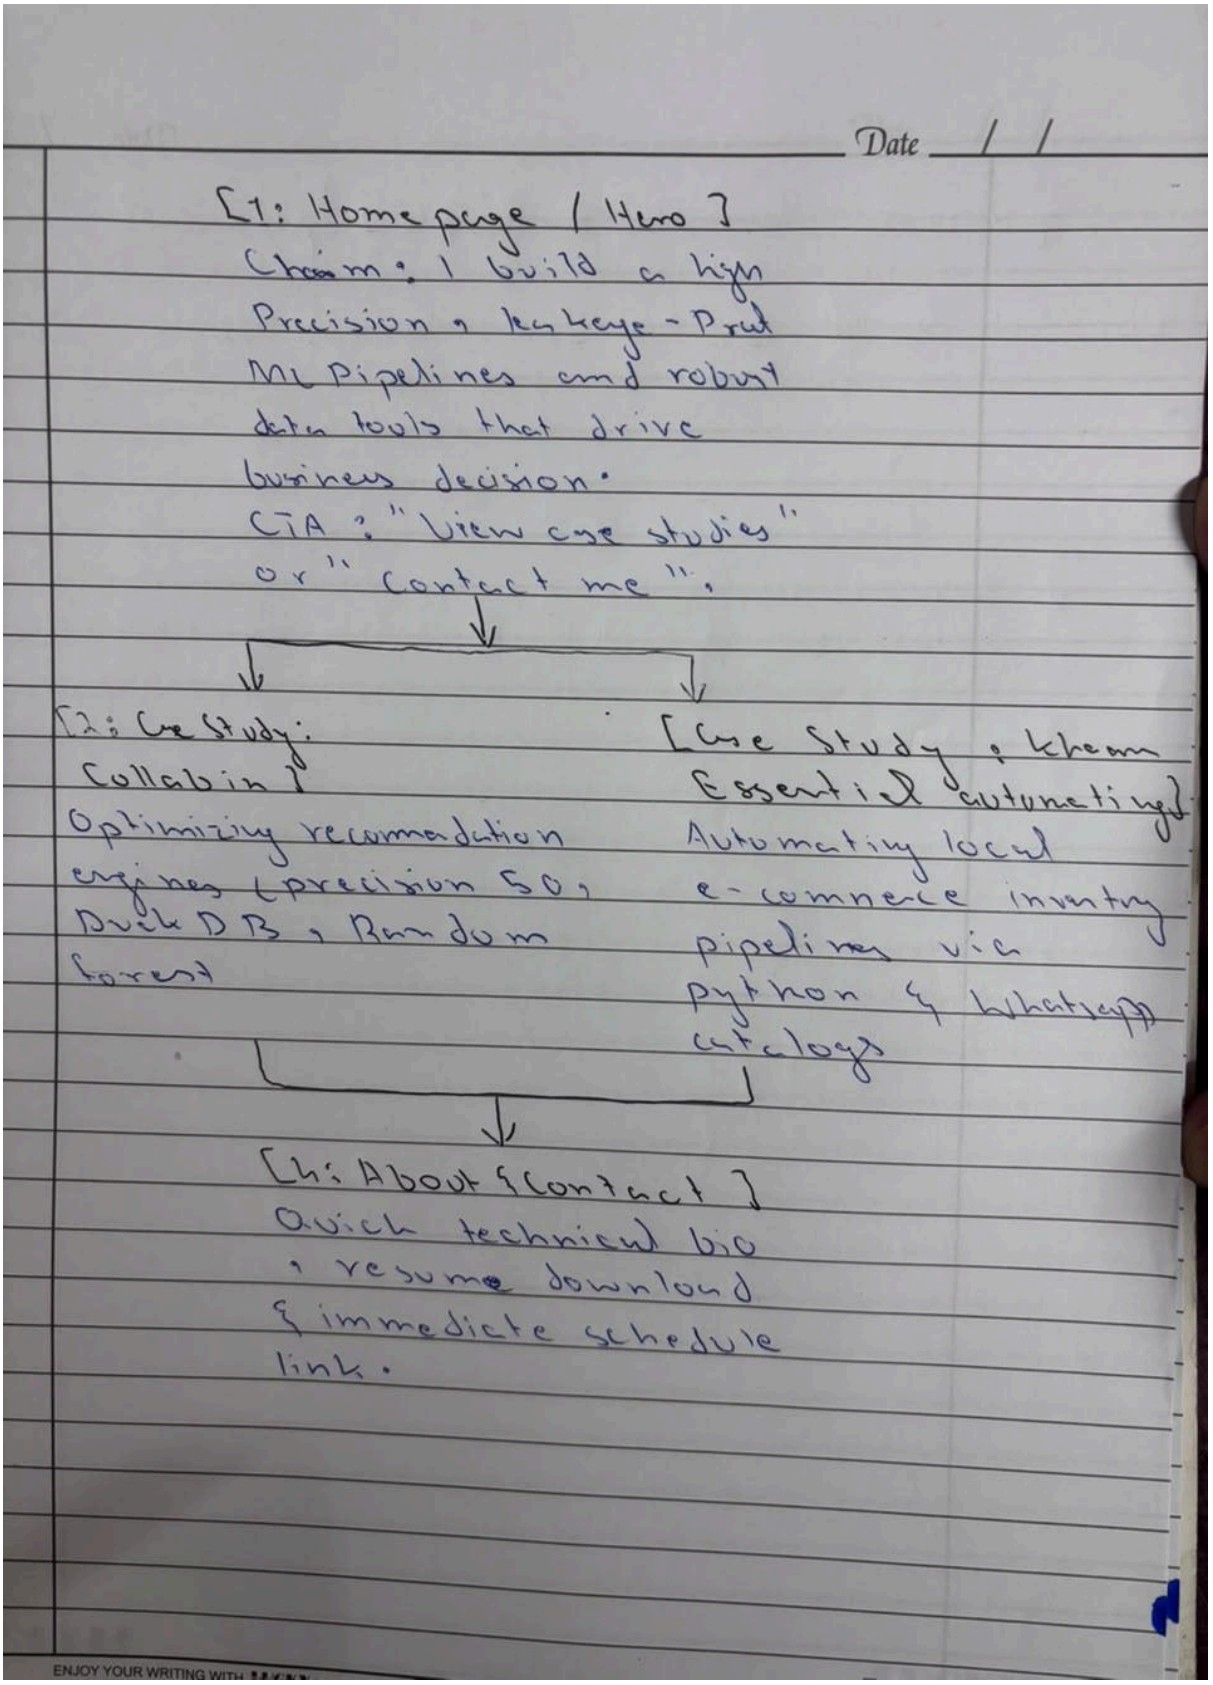

Case Study 2 paragraph: signed off. That's done — move it to final copy.

The homepage hero has one real problem, and it's the first line: **"turn complex raw metrics into defensible business outcomes."**

Read that sentence back to yourself with the same eye we've used on every stat for the last four rounds. "Complex raw metrics" and "defensible business outcomes" are content-free phrases — they're the exact register as "decision-support tools" and "in real time," the filler we've cut twice already. They sound like a sentence, they fill space, but if you deleted them the page would lose nothing. That's the fluff test: does removing it cost you information? Here, no.

Bigger issue: you dropped your sharpest word.
Your original proof statement had "leakage-proof." That word is a filter — it's the one phrase on this entire site that a technical recruiter or hiring manager who actually knows ML will read and immediately clock as "this person understands validation methodology, not just sklearn.fit()." Most junior candidates don't know what leakage even is, let alone that it's the thing that silently invalidates half of Kaggle-tier portfolios. "High-precision" is generic — every bootcamp grad's LinkedIn says "high-precision." "Leakage-proof" is a credential. Don't lose it in a rewrite for flow.

Fix:

> **Hero Title:**
> "I build leakage-proof ML pipelines and high-throughput data tools that hold up under real validation and real load."

(Note: "hold up under real validation and real load" directly nods to both case studies — client-grouped validation for CS1, the 70x/4.2M-row load for CS2 — without needing an abstract noun like "outcomes.")

Sub-headline: this one's actually good, minor tightening only.
It's dense with your two hardest numbers, which is exactly right for a sub-headline's job. One small fix: "over millions of rows" — you have the exact figure, 4.2 million. Use it. "Millions" is the same rounding-up instinct we killed on "84.3 seconds → 70x." You paid for the precise number, spend it:

> "From client-grouped validation pipelines that hold 84% Precision@50 with zero target leakage, to DuckDB query optimization delivering a 70x speedup across 4.2M rows — I build software engineers trust and businesses can scale on."

CTA: fine as is. "View My Systems & Case Studies" is clean, directional, no fluff. Don't touch it.

One last check before you call the whole 8-week build done: read the homepage, then Case Study 1, then Case Study 2, then About, back to back, out loud. If any sentence in that sequence could be deleted without losing a number, a tool name, or a decision you made — cut it. That's your standing rule now, not just for this project.

Task 1: Refactoring & Optimizing Recommendation Algorithms

- **What "Done Well" Means:**
    - **Execution Speed:** The optimized DuckDB/Pandas Python script must execute over a target dataset within defined latency bounds.
    - **Rigorous Data Separation:** The pipeline must feature strict index partitioning to ensure zero target or features leakage between training splits.
    - **Maintainability:** Code must score 10/10 on structural linting (Flake8/Pylint) and contain descriptive docstrings for team handoffs.

Task 2: Writing Copy for Khaam Essentials

- **What "Done Well" Means:**
    - **Brand Alignment:** Avoids robotic, overly generic AI buzzwords (like *"Revolutionize your hair game!"* or *"Delve into"*).
    - **Authenticity:** Masters a natural, storytelling tone that explicitly highlights our core brand pillars: organic, chemical-free, and lovingly handcrafted by our mother.
    - **Formatting Constraints:** Copy must adhere to strict platform boundaries (e.g., under 150 characters for Instagram Bios or structured bullet points for WhatsApp catalogs) and naturally integrate target keywords.

Task 3: Building Engineering Course Study & Problem Guides

- **What "Done Well" Means:**
    - **Mathematical Accuracy:** Generated practice problems for numerical analysis (e.g., root-finding, interpolation) or instrumentation process control must contain 100% correct, verifiable step-by-step solutions.
    - **Syllabus Mapping:** The study guides must map perfectly to my university course handbook chapters without introducing irrelevant, out-of-scope concepts.
    - **Clarity:** Formulas must be cleanly formatted using clear Markdown tables and standard notation.

What Are You Proving?

1. The One-Paragraph Proof Statement

> I build leakage-proof machine learning pipelines. I am proving this to a Lead Machine Learning Engineer or Engineering Manager hiring for an early-stage ML team who needs developers capable of executing reliable validation structures without hand-holding. The single, high-leverage action I want them to take after landing on my page is to click my scheduler link and book a 30-minute technical interview call. I prove my claim by showing them how I designed a strict client-grouped split that prevents target leakage and holds an 84% Precision@50 on entirely unseen domains.

2. The One-Line "Why"

This portfolio needs to exist because a standard resume or LinkedIn profile only lists tools and courses; it cannot prove that I possess the engineering judgment to design mathematically rigorous, leakage-free validation splits that prevent catastrophic model decay in production.

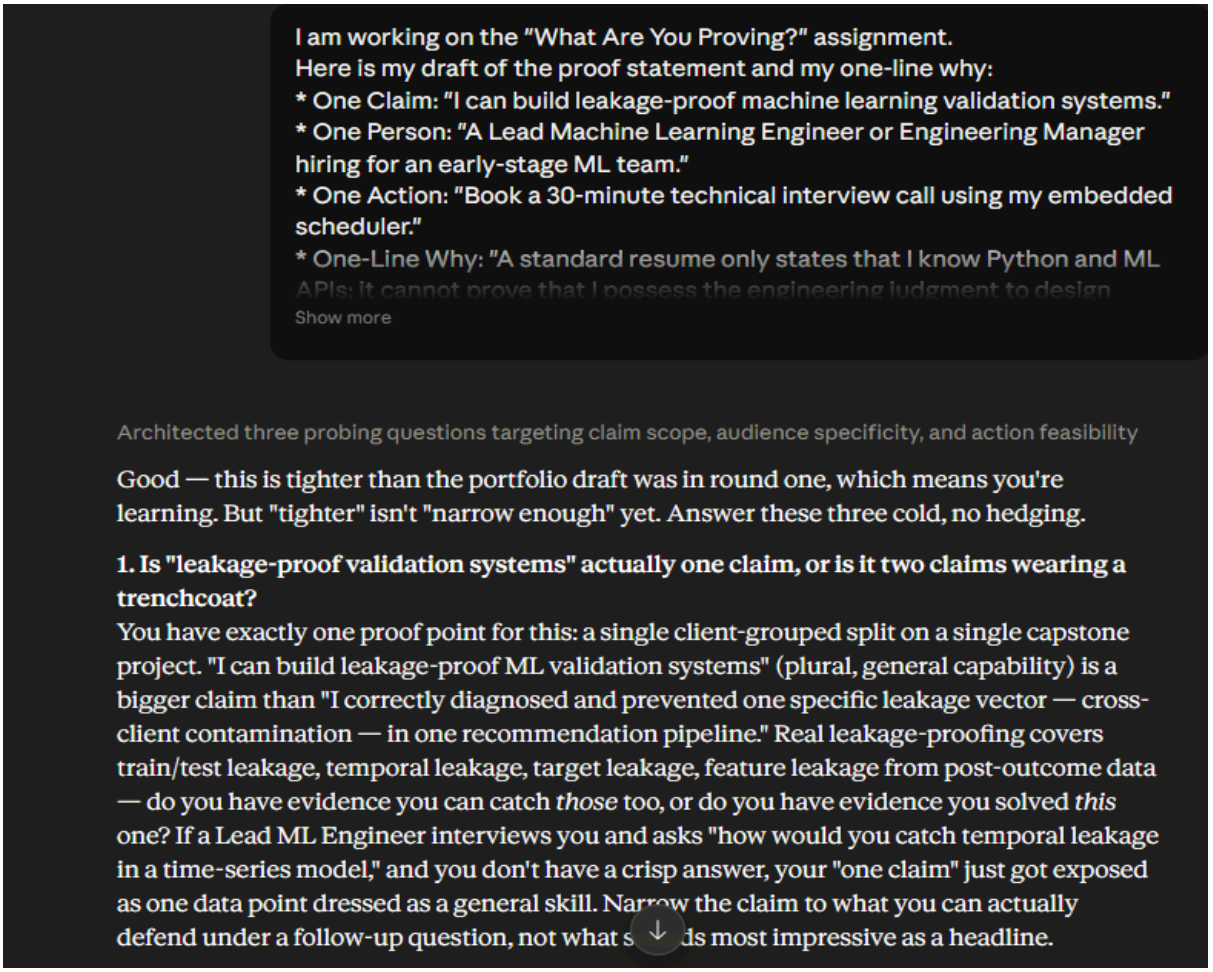

2. "Lead ML Engineer or Engineering Manager hiring for an early-stage ML team" — pick one, because they don't want the same proof.
A Lead ML Engineer evaluating you is checking your technical judgment — can you actually reason about bias-variance tradeoffs, split design, metric selection. An Engineering Manager at an early-stage team is checking something different — can you ship without hand-holding, will you slow the team down, are you low-maintenance. Those are different audiences who read the same case study looking for different signals. Your "one person" isn't one person, it's a category with two different hot buttons. If you write your pitch for both, you'll write it for neither with real precision. Which one are you actually optimizing this proof for? Say the specific job title you'd search on LinkedIn right now, not the role family.

3. "Book a 30-minute technical interview call using my embedded scheduler" — realistic, or wishful?
Walk me through the actual mechanism. A senior ML lead is not browsing student portfolio sites and spontaneously booking calendar slots with strangers off a scheduler widget — that's not how technical hiring conversations start, for anyone, at any level. What's the actual path this person takes to land on your site in the first place — are you sending it to them directly (cold outreach, referral, application follow-up), or are you hoping for organic discovery? If it's the former, your "one action" is really "reply to my message and agree to a call," and the portfolio's job is to survive scrutiny *after* that reply, not to generate the first contact. If you're expecting the scheduler to do outreach's job, that's the same rounding-up mistake we killed in your case study numbers — dressing up a hopeful outcome as a designed mechanism. Be honest about what stage of the funnel this asset is actually built for.

Answer all three before you finalize the assignment. Don't soften them on the way back to me.

You just completely demystified the hiring funnel. Let's strip out the wishful thinking and narrow this down to a razor-sharp, defensible baseline.
Here are my direct answers:
1. One Claim vs. Two in a Trenchcoat
It is absolutely one specific project wearing a general capability's trenchcoat. I am narrowing the claim from "leakage-proof validation systems" to:
"I design leak-free validation splits for client-grouped and tabular datasets."
This aligns exactly with my proof point (unseen client-domain evaluation) and is a specific, defensible skill. If an interviewer asks about time-series or
Show more

3:07 PM

Close, but not quite. Two things left, and one of them is a real problem, not a polish issue.

1. "Prevent catastrophic model decay in production" — you're back to rounding up.
Stop and check this against what you actually did. You ran a client-grouped split on a capstone project and measured 84% Precision@50 on held-out domains. That's a well-designed *offline validation methodology*. You did not deploy a model to production, monitor it over time, and prevent decay — decay is a live, post-deployment phenomenon (data drift, concept drift, degrading predictions over weeks/months). Claiming your validation split "prevents catastrophic model decay in production" is claiming a downstream outcome you have zero evidence for. If the Lead ML Engineer asks "so you monitored this in prod and saw decay rates improve?" — you don't have an answer, because it never went to production. This is the exact same move as "50% of budgets" and "memory overloads" — a true, specific, defensible finding (leak-free offline validation) stretched into a bigger, unverifiable claim (production decay prevention) because the bigger claim sounds more consequential. Cut it. Say what you actually proved: that your split design gives an honest, generalizable estimate of real-world performance *before* deployment — which is itself the whole point of leakage-proofing, you don't need to reach further than that.

2. The "why" still has one soft phrase: "skeptical Lead ML Engineer."
Minor, but "skeptical" is doing decorative work — it's an adjective describing the reader's mood, not a fact about your proof. Cut it or replace with something concrete about what they're actually skeptical *of* — e.g., "a Lead ML Engineer who has seen enough portfolios claiming strong metrics on leaky pipelines to distrust the number until the split methodology is shown." That's longer, but it's specific and it earns the word "skeptical" instead of asserting it.

Everything else stands: the narrowed claim, the single persona, the honest funnel position (portfolio validates the pitch, doesn't generate it) — all of that is exactly right and matches your evidence precisely. That's the actual skill you've been building over this whole exercise: catching your own claim the moment it exceeds your evidence, before someone else catches it for you.

Fix the "production decay" line and this is done. Send me the one corrected sentence, not the whole paragraph again — I want to see you make just the one surgical cut.