

Section A: Personal Workflow Audit


#	Recurring Weekly Task	Classification	One-Line Rationale
1	Handcrafting organic shampoo batches	Just Me	Physical, artisan-level compounding done in small batches by my mother that requires tactile quality control.
2	Negotiating raw ingredient prices with local vendors	Just Me	Requires building personal trust, navigating local wholesale markets, and real-time human negotiation.
3	Testing and debugging physical hardware circuits	Just Me	Requires physical laboratory instrumentation, multimeters, and manual wiring that AI cannot physically interact with.
4	Writing complex SQL & DuckDB queries	Collaborate with AI	I provide the database schema and query intent; the AI rapidly drafts the initial syntax, which I then optimize for performance.
5	Drafting social copy & catalog descriptions for @Khaam.essentials	Collaborate with AI	I provide the product benefits; the AI refines the tone and structures the copy, which I then localize for our target audience.
6	Translating numerical analysis algorithms into Python scripts	Collaborate with AI	Great for translating step-by-step mathematical logic (like Jacobi or Gauss-Seidel) into structured, clean code.

7	Synthesizing engineering handbook chapters into review sheets	Delegate to AI with Review	I upload the raw PDF chapter; the AI extracts key formulas and definitions, which I audit against my lecture notes.
8	Reviewing documentation on rule-based recommendation scoring	Delegate to AI with Review	The AI summarizes dense API/framework docs so I can quickly extract implementation patterns for my internship tasks.
9	Triaging basic customer FAQs on Instagram & WhatsApp	Delegate to AI with Review	Standard questions about shipping, product availability, or price are drafted by AI templates, requiring only my manual click to send.
10	Formatting, linting, and cleaning local Python code bases	Fully Automate	Easily managed using automated Git hooks, Black, and Flake8 formatters running locally on file save.
11	Syncing raw inventory lists with spreadsheet stock counts	Fully Automate	Handled by simple Google Sheets script triggers that auto-update stock levels whenever a WhatsApp order form is logged.
12	Generating mock technical interview questions	Just Me	Helps me practice core engineering and discrete modeling concepts through interactive chat sessions.

Section B: Setup Evidence & Custom Instructions

Paste this description into your document, and then **insert your screenshot** directly below it:

- **Claude.ai / ChatGPT Free Accounts:** Confirmed active.
- **Anthropic Academy:** Enrolled in *Claude 101* (Module 1 completed).




Instructions for Claude

Claude will keep these in mind across chats and Cowork within [Anthropic's guidelines](#). [Learn more](#)

Professional Summary

Technical and detail-oriented Engineering and Computer Science student with proven experience in backend optimization, numerical simulation, and data-driven systems. Currently leveraging Python, SQL, and DuckDB to develop high-performance recommendation engines and automated workflows. Strong background in systematic validation, functional testing, and rigorous analytical reasoning.

Section C: Selected Tasks & Success Metrics



Case Study 2 paragraph: signed off. That's done — move it to final copy.

The homepage hero has one real problem, and it's the first line: **"turn complex raw metrics into defensible business outcomes."**

Read that sentence back to yourself with the same eye we've used on every stat for the last four rounds. "Complex raw metrics" and "defensible business outcomes" are content-free phrases — they're the exact register as "decision-support tools" and "in real time," the filler we've cut twice already. They sound like a sentence, they fill space, but if you deleted them the page would lose nothing. That's the fluff test: does removing it cost you information? Here, no.

Bigger issue: you dropped your sharpest word.

Your original proof statement had "leakage-proof." That word is a filter — it's the one phrase on this entire site that a technical recruiter or hiring manager who actually knows ML will read and immediately clock as "this person understands validation methodology, not just `sklearn.fit()`." Most junior candidates don't know what leakage even is, let alone that it's the thing that silently invalidates half of Kaggle-tier portfolios. "High-precision" is generic — every bootcamp grad's LinkedIn says "high-precision." "Leakage-proof" is a credential. Don't lose it in a rewrite for flow.

Fix:

Hero Title:

"I build leakage-proof ML pipelines and high-throughput data tools that hold up under real validation and real load."

(Note: "hold up under real validation and real load" directly nods to both case studies — client-grouped validation for CS1, the 70x/4.2M-row load for CS2 — without needing an

(Note: "hold up under real validation and real load" directly nods to both case studies — client-grouped validation for CS1, the 70x/4.2M-row load for CS2 — without needing an abstract noun like "outcomes.")

Sub-headline: this one's actually good, minor tightening only.

It's dense with your two hardest numbers, which is exactly right for a sub-headline's job. One small fix: "over millions of rows" — you have the exact figure, 4.2 million. Use it. "Millions" is the same rounding-up instinct we killed on "84.3 seconds → 70x." You paid for the precise number, spend it:

"From client-grouped validation pipelines that hold 84% Precision@50 with zero target leakage, to DuckDB query optimization delivering a 70x speedup across 4.2M rows — I build software engineers trust and businesses can scale on."

CTA: fine as is. "View My Systems & Case Studies" is clean, directional, no fluff. Don't touch it.

One last check before you call the whole 8-week build done: read the homepage, then Case Study 1, then Case Study 2, then About, back to back, out loud. If any sentence in that sequence could be deleted without losing a number, a tool name, or a decision you made — cut it. That's your standing rule now, not just for this project.

📄 🔊 📄 📄 📄

Task 1: Refactoring & Optimizing Recommendation Algorithms

- **What "Done Well" Means:**
 - **Execution Speed:** The optimized DuckDB/Pandas Python script must execute over a target dataset within defined latency bounds.
 - **Rigorous Data Separation:** The pipeline must feature strict index partitioning to ensure zero target or features leakage between training splits.
 - **Maintainability:** Code must score 10/10 on structural linting (Flake8/Pylint) and contain descriptive docstrings for team handoffs.

Task 2: Writing Copy for Khaam Essentials

- **What "Done Well" Means:**
 - **Brand Alignment:** Avoids robotic, overly generic AI buzzwords (like *"Revolutionize your hair game!"* or *"Delve into"*).
 - **Authenticity:** Masters a natural, storytelling tone that explicitly highlights our core brand pillars: organic, chemical-free, and lovingly handcrafted by our mother.
 - **Formatting Constraints:** Copy must adhere to strict platform boundaries (e.g., under 150 characters for Instagram Bios or structured bullet points for WhatsApp catalogs) and naturally integrate target keywords.

Task 3: Building Engineering Course Study & Problem Guides

- **What "Done Well" Means:**
 - **Mathematical Accuracy:** Generated practice problems for numerical analysis (e.g., root-finding, interpolation) or instrumentation process control must contain 100% correct, verifiable step-by-step solutions.
 - **Syllabus Mapping:** The study guides must map perfectly to my university course handbook chapters without introducing irrelevant, out-of-scope concepts.
 - **Clarity:** Formulas must be cleanly formatted using clear Markdown tables and standard notation.